

Data Wrangling Part 1

Cecilia Y. Sui and all other TAs

Contents

1	Tips on Writing Code	1
1.1	Coding Style	1
1.2	Organizing Your Files	3
1.3	In-class exercises:	4
2	Loading and Saving Data	4
2.1	Built-in Datasets	4
2.2	Loading Data from Sources	5
2.3	Saving Data to Files	6
2.4	In-class exercises:	7
3	Data Wrangling Part 1	7
3.1	Selecting Values	7
3.2	In-class exercises:	12
3.3	Dollar Signs and Double Brackets	15
4	Modifying Values	18
4.1	Changing Values in Place	18
4.2	Logical Subsetting	20
4.3	In-class exercises:	22
4.4	In-class exercises:	24
4.5	Missing Information (NA)	26
4.6	In-class exercises:	28

1 Tips on Writing Code

By default, RStudio will save your current workspace when you quit. While convenient, this can mean that you make one-off changes to your data and forget to save that command in your script. Starting with a fresh session every time you open RStudio means you'll learn to keep every step of your analysis in your script, and you'll know that you can get back to where you were by rerunning the script.

1.1 Coding Style

There are two very good reasons to try to write your code in a clear, understandable way:

- Other people might need to use your code to replicate your studies.
- You might need to use your code, a few weeks/months/years after you've written it.

It's possible to write R code that "works" perfectly, and produces all the results and output you want, but proves very difficult to make changes to when you have to come back to it (because a reviewer asked for one additional analysis, etc. which happens very often.)

1.1.1 Formatting Tips

You can improve the readability of your code a lot by following a few simple rules:

- Put spaces between and around variable names and operators (`=+-*/`)
- Break up long lines of code
- Use meaningful variable names composed of 2 or 3 words” (avoid abbreviations unless they’re very common and you use them very consistently)
- Indentation (2 or 4 spaces)

These rules can mean the difference between this:

```
lm1=lm(y~grp+grpTime,mydf,subset=sext1== "m")
```

and this:

```
male_difference = lm(DepressionScore ~ Group + GroupTimeInteraction,
                      data = interview_data,
                      subset = BaselineSex == "Male")
```

R will treat both pieces of code exactly the same, but for any humans reading, the nicer layout and meaningful names make it much easier to understand what’s happening, and spot any errors in syntax or intent.

1.1.2 Keep a Consistent Style

Try to follow a consistent style for naming things, e.g. using **** snake_case for all your variable names in your R code, and TitleCase **** for the columns in your data. Either style is probably better than lowercase with no spacing allmashedtogether. DO NOT use . in names (`do.not.do.this`). It doesn’t particularly matter what that style is, as long as you’re consistent.

1.1.3 Write LOTS of Comments

As any other code, one of the best things you can do to make R code readable and understandable is write comments. R ignores lines that start with # so you can write whatever you want and it won’t affect the way your code runs.

Comments that explain why something was done are great:

```
# reverse the score for question 3
data$DepressionQ3 = 4 - data$DepressionQ3
```

Comments that explain what is being done are also helpful. These comments might appear redundant for people who already understand R code, but can be super helpful for you to learn what is going on and for grading purposes as well.

```
# Calculate the mean of the anxiety scores
anxiety_mean = mean(data$AnxietyTotal)
```

You can also put comments that remind you to fix certain errors.

```
##### FIX HERE: This fails to converge #####
```

1.1.4 Google’s R Style Guide

In QPM I, we would like you to follow the R Programming Style Guide by Google to make your R code easier to read, share, and verify. For more details, please check out this link: <https://google.github.io/styleguide/Rguide.html>

1.2 Organizing Your Files

Working directories are extremely useful for keeping your work organized. When you have hundreds of R files along with images and data, it is much easier to find things using directories instead of putting all the files in one folder. A directory is basically a folder that you have created to save all your work.

1.2.1 Get the working directory

“Each time you open R, it links itself to a directory on your computer, which R calls the working directory. This is where R will look for files when you attempt to load them, and it is where R will save files when you save them. The location of your working directory will vary on different computers. You can find out which directory by running the `getwd()` function, which stands for “get the working directory”.

By default your working directory is set to be *the location of your current file*. For example, if you have created an R script as `~/Documents/Rcamp2026/test.R`, the working directory for the R script will be `~/Documents/Rcamp2026/`. Same for R markdown files.

```
getwd()
```

```
## [1] "/Users/chanhuk/Dropbox/courses/RCamp2026/Day2/Lecture"
```

1.2.2 Change the working directory

You can place data files straight into the folder that is your working directory, or you can move your working directory to where your data files are. You can move your working directory to any folder on your computer with the function `setwd()`. Just give `setwd()` the file path to your new working directory. This may be from the root directory (starting with `/` on a Mac), it may include a double-dot `(..)` to move locally up a folder from the current directory, and it may include a path from the current directory. You can set the working directory to a folder dedicated to whichever project you are currently working on. That way you can keep all of my data, scripts, graphs, and reports in the same place. For example:

```
setwd("/Users/river/Documents/GitHub/RCamp2023/Day2/Lecture")
```

If the file path does not begin with your root directory, R will assume that it begins at your current working directory.

CAUTION When you are working on Rmarkdown, `setwd()` only **temporarily** changes working directory for *a code chunk*. If you want to set the working directory for the *whole R markdown document*, include or change `knitr::opts_knit$set(root.dir = '{PATH}')`. However, the best practice is just to create R script / R markdown file on the working directory and do not change it.

1.2.3 List files in the working directory

You can see what files are in your working directory with `list.files()`. If you see the file that you would like to open in your working directory, then you are ready to proceed. How you open files in your working directory will depend on which type of file you would like to open.

```
list.files()
```

```
## [1] "booleans.png"      "data"          "dataset.csv"       "Day2_Lec2022.pdf"
## [5] "Day2_Lec2022.Rmd"   "logicals.png"    "school_loc.csv"    "stuff.RData"
## [9] "stuff.RDS"
```

tips If you ever worked with UNIX terminal, this reminds you of `ls` command. R do have `ls()` command too, but that command will print out the list of variables.

1.2.4 Create a new directory

To create a directory in R, use the `dir.create(path)` method. The `dir.create()` method accepts a folder generated in the current working directory or specifies a path. Another way is to go to Session > Set Working

Directory.

```
# This creates a new directory inside the current directory at path = /Users/river/Documents/GitHub/RCa  
dir.create("./pset2")
```

1.3 In-class exercises:

1. Get your current working directory.

2 Loading and Saving Data

In Day 1's lecture, we covered how to build a dataframe by hand-typing all the values in each row and column. In this section, we will explore the built-in datasets provided by base R. More importantly, we will cover how to load data from sources like csv files or excel sheets, and how to save the new dataframe that you have made modifications to.

2.1 Built-in Datasets

R comes with many data sets preloaded in the datasets package, which comes with base R. These data sets are not very interesting, but they give you a chance to test code or make a point without having to load a data set from outside R. You can see a list of R's data sets as well as a short description of each by running:

```
help(package = "datasets")
```

To use a data set, just type its name. Each data set is already presaved as an R object. For example:

```
iris
```

```
##   Sepal.Length Sepal.Width Petal.Length Petal.Width   Species  
## 1         5.1       3.5      1.4       0.2    setosa  
## 2         4.9       3.0      1.4       0.2    setosa  
## 3         4.7       3.2      1.3       0.2    setosa  
## 4         4.6       3.1      1.5       0.2    setosa  
## 5         5.0       3.6      1.4       0.2    setosa  
## 6         5.4       3.9      1.7       0.4    setosa  
## 7         4.6       3.4      1.4       0.3    setosa  
## 8         5.0       3.4      1.5       0.2    setosa  
## 9         4.4       2.9      1.4       0.2    setosa  
...  
mtcars
```

```
##          mpg cyl  disp  hp drat    wt  qsec vs am gear carb  
## Mazda RX4     21.0   6 160.0 110 3.90 2.620 16.46  0  1    4    4  
## Mazda RX4 Wag  21.0   6 160.0 110 3.90 2.875 17.02  0  1    4    4  
## Datsun 710    22.8   4 108.0  93 3.85 2.320 18.61  1  1    4    1  
## Hornet 4 Drive 21.4   6 258.0 110 3.08 3.215 19.44  1  0    3    1  
## Hornet Sportabout 18.7   8 360.0 175 3.15 3.440 17.02  0  0    3    2  
## Valiant      18.1   6 225.0 105 2.76 3.460 20.22  1  0    3    1  
## Duster 360    14.3   8 360.0 245 3.21 3.570 15.84  0  0    3    4  
## Merc 240D     24.4   4 146.7  62 3.69 3.190 20.00  1  0    4    2  
## Merc 230      22.8   4 140.8  95 3.92 3.150 22.90  1  0    4    2
```

...

You can use `str()` function to view the structure of the data and `summary()` function to summarize the dataset. The function tells you the minimum, maximum and quartiles for each numeric column. If there are NAs in your dataframe, it also tells you how many NAs there are in each column.

```
str(iris)

## 'data.frame': 150 obs. of 5 variables:
## $ Sepal.Length: num 5.1 4.9 4.7 4.6 5 5.4 4.6 5 4.4 4.9 ...
## $ Sepal.Width : num 3.5 3 3.2 3.1 3.6 3.9 3.4 3.4 2.9 3.1 ...
## $ Petal.Length: num 1.4 1.4 1.3 1.5 1.4 1.7 1.4 1.5 1.4 1.5 ...
## $ Petal.Width : num 0.2 0.2 0.2 0.2 0.2 0.4 0.3 0.2 0.2 0.1 ...
## $ Species     : Factor w/ 3 levels "setosa","versicolor",...: 1 1 1 1 1 1 1 1 1 1 ...

summary(iris)

##   Sepal.Length   Sepal.Width   Petal.Length   Petal.Width
## Min.   :4.300   Min.   :2.000   Min.   :1.000   Min.   :0.100
## 1st Qu.:5.100  1st Qu.:2.800  1st Qu.:1.600  1st Qu.:0.300
## Median :5.800  Median :3.000  Median :4.350  Median :1.300
## Mean   :5.843  Mean   :3.057  Mean   :3.758  Mean   :1.199
## 3rd Qu.:6.400  3rd Qu.:3.300  3rd Qu.:5.100  3rd Qu.:1.800
## Max.   :7.900  Max.   :4.400  Max.   :6.900  Max.   :2.500

##   Species
## setosa    :50
## versicolor:50
## virginica:50
##
##
```

However, when doing your own research, you will need to load your own data into R from a wide variety of file formats. Before you can load any data files into R, you will need to tell R where your working directory is.

2.2 Loading Data from Sources

2.2.1 CSV Files

R provides several read functions for you to load data from files, which are shortcuts for `read.table` with different default arguments. The most commonly used one is `read.csv()`. You will probably use this function hundreds of times during our sequence of methods courses here at WashU. As suggested by its name, the function `read.csv` allows you to load in comma-separated-variable (CSV) files. For example:

```
school_loc <- read.csv("school_loc.csv")
# remove(school_loc)
```

2.2.2 HTML Links

Many data files are made available on the Internet at their own web address. If you are connected to the Internet, you can open these files straight into R with `read.csv()` (or other functions based on `read.table()`). You can pass a web address into the file name argument for any of R's data-reading functions. For example:

```
school_loc <- read.csv("https://data-nces.opendata.arcgis.com/datasets/a15e8731a17a46aabc452ea607f172c0")
# View(school_loc)
# https://www.data.gov/
```

Just make sure that the web address links directly to the file and not to a web page that links to the file.

Usually, when you visit a data file's web address, the file will begin to download or the raw data will appear in your browser window.

2.2.3 Native R Data Formats

R provides two file formats of its own for storing data, `.RDS` and `.RData`. These are known to take up smaller spaces than csvs and faster to read and write. RDS files can store a *single R object*, and RData files can store *multiple R objects*.

You can open a RDS file with `readRDS`. For example, if the royal flush data was saved as `poker.RDS`, you could open it with:

```
school_loc <- readRDS("stuff.RDS")
```

Opening RData files is even easier. Simply run the function `load` with the file:

```
# be careful not to assign overlapping data name
# a <- 'text'
load("stuff.RData")
```

There's no need to assign the output to an object. The R objects in your RData file will be loaded into your R session *with their original names*. RData files can contain multiple R objects, so loading one may read in multiple objects. `load()` doesn't tell you how many objects it is reading in, nor what their names are, so it pays to know a little about the RData file before you load it.

If worse comes to worst, you can keep an eye on the environment pane in RStudio as you load an RData file. It displays all of the objects that you have created or loaded during your R session. Another useful trick is to put parentheses around your `load` command like so, `(load("file.RData"))`. This will cause R to print out the names of each object it loads from the file.

Both `readRDS()` and `load()` take a file path as their first argument, just like R's other read and write functions. If your file is in your working directory, the file path will be the file name.

2.3 Saving Data to Files

2.3.1 Native R Data Formats

You can save an R object like a data frame as either an RData file or an RDS file. RData files can store multiple R objects at once, but RDS files are the better choice because they foster reproducible code.

To save data as an RData object, use `save()` function. To save data as a RDS object, use `saveRDS()` function. In each case, the first argument should be the name of the R object you wish to save. You should then include a file argument that has the file name or file path you want to save the data set to. Remember to put file extensions!

For example, if you have three R objects, `a`, `b`, and `c`, you could save them all in the same RData file and then reload them in another R session:

```
a <- 1
b <- 2
c <- 3
save(a, b, c, file = "stuff.RData")
load("stuff.RData")
```

However, if you forget the names of your objects or give your file to someone else to use, it will be difficult to determine what was in the file—even after you (or they) load it. The user interface for RDS files is much more clear. You can save only one object per file, and whoever loads it can decide what they want to call their new data. As a bonus, you don't have to worry about load overwriting any R objects that happened to have the same name as the objects you are loading:

```
saveRDS(a, file = "stuff.RDS") # saving a to RDS file
d <- readRDS("stuff.RDS")
# readRDS("stuff.RDS")
```

Saving your data as an R file offers some advantages over saving your data as a plain-text file. R automatically compresses the file and will also save any R-related metadata associated with your object. This can be handy if your data contains factors, dates and times, or class attributes. You won't have to reparse this information into R the way you would if you converted everything to a text file.

On the other hand, R files cannot be read by many other programs, which makes them inefficient for sharing. They may also create a problem for long-term storage if you don't think you'll have a copy of R when you reopen the files.

2.3.2 CSV Files

You can also use `write.csv()` to save your data as a .csv file. The first argument is the R object that contains your data set. The file argument is the file name (including extension) that you wish to give the saved data. By default, the function will save your data into your working directory. However, you can supply a file path to the file argument. R will oblige by saving the file at the end of the file path. If the file path does not begin with your root directory, R will append it to the end of the file path that leads to your working directory.

For example, you can save the school location data frame to a subdirectory named data within your working directory with the command:

```
# write.csv(r_object, file = filepath, row.names = FALSE)
write.csv(school_loc, "./data/school_loc.csv", row.names = FALSE)
```

Keep in mind that `write.csv()` cannot create new directories on your computer. Each folder in the file path must exist before you try to save a file with it.

The `row.names()` argument prevents R from saving the data frame's row names as a column in the plain-text file. You might have noticed that R automatically names each row in a data frame with a number.

2.4 In-class exercises:

1. Load the world total fertility rate dataset from `dataset.csv`. Name the dataframe `TFR`.
2. Create a new directory named `data`.
3. Save the dataframe as “`newdata.csv`” into the new directory that you just created.

3 Data Wrangling Part 1

In Day 1's lecture, we covered how to create a vector, a list, and most importantly a data frame. Now we need to know how to manipulate the individual values inside your data frame. You can select values within an R object with R's notation system.

3.1 Selecting Values

R has a notation system that lets you extract values from R objects. To extract a value or set of values from a data frame, write the data frame's name followed by a pair of hard brackets:

Between the brackets will go two indices separated by a comma. The indices tell R which values to return. R will use the first index to subset the rows of the data frame and the second index to subset the columns.

```

# View(school_loc)
school_loc[ , ]

# str(school_loc)
# summary(school_loc)

# First few rows
head(school_loc)

# Last few rows
tail(school_loc)

# typeof(school_loc)
# class(school_loc)

```

You have a choice when it comes to writing indices. There are six different ways to write an index for R, and each does something slightly different. They are all very simple and quite handy, so let's take a look at each of them. You can create indices with:

- * Positive integers
- * Negative integers
- * Zero
- * Blank spaces
- * Logical values
- * Names

The simplest and most commonly used of these is positive integers.

3.1.1 Positive integers

R treats positive integers just like the `ij` notation used in linear algebra: `school_loc[i,j]` will return the value of dataframe that is in the `i`th row and the `j`th column. Notice that `i` and `j` only need to be integers in the mathematical sense.

```
head(school_loc)
```

	X	Y	OBJECTID	UNITID	NAME
## 1	-86.56850	34.78337	1	100654	Alabama A & M University
## 2	-86.79935	33.50570	2	100663	University of Alabama at Birmingham
## 3	-86.17401	32.36261	3	100690	Amridge University
## 4	-86.64045	34.72456	4	100706	University of Alabama in Huntsville
## 5	-86.29568	32.36432	5	100724	Alabama State University
## 6	-87.52959	33.20702	6	100733	University of Alabama System Office
##			STREET	CITY STATE	ZIP STFIP CNTY
## 1	4900 Meridian Street		Normal	AL	35762 01 01089
## 2	Administration Bldg Suite 1070	Birmingham		AL	35294-0110 01 01073
## 3	1200 Taylor Rd	Montgomery		AL	36117-3553 01 01101
## 4	301 Sparkman Dr	Huntsville		AL	35899 01 01089
## 5	915 S Jackson Street	Montgomery		AL	36104-0271 01 01101
## 6	500 University Blvd. East	Tuscaloosa		AL	35401 01 01125
##	NMCNTY	LOCALE	LAT	LON	CBSA NMCBSA
## 1	Madison County		12	34.78337	-86.56850 26620 Huntsville, AL
## 2	Jefferson County		12	33.50570	-86.79935 13820 Birmingham-Hoover, AL
## 3	Montgomery County		12	32.36261	-86.17401 33860 Montgomery, AL
## 4	Madison County		12	34.72456	-86.64045 26620 Huntsville, AL
## 5	Montgomery County		12	32.36432	-86.29568 33860 Montgomery, AL
## 6	Tuscaloosa County		12	33.20701	-87.52959 46220 Tuscaloosa, AL
##	CBSATYPE	CSA			NMCSA NECTA NMNECTA CD SLDL
## 1	1	290			Huntsville-Decatur, AL N N 0105 01019
## 2	1	142			Birmingham-Hoover-Talladega, AL N N 0107 01055
## 3	1	388			Montgomery-Selma-Alexander City, AL N N 0102 01074

```

## 4      1 290      Huntsville-Decatur, AL      N      N 0105 01006
## 5      1 388 Montgomery-Selma-Alexander City, AL      N      N 0107 01077
## 6      1   N      N      N 0107 01063
##     SLDU SCHOOLYEAR
## 1 01007 2020-2021
## 2 01018 2020-2021
## 3 01025 2020-2021
## 4 01002 2020-2021
## 5 01026 2020-2021
## 6 01021 2020-2021

# Extract value at the 1st row and the 2nd column
school_loc[1,2]

```

```
## [1] 34.78337
```

Notice that the row and column indices for R starts at 1, instead of 0 in some other programming languages. This is a more intuitive way of indexing things, but can be confusing if you go back and forth.

To extract more than one value, use *a vector* of positive integers. For example, you can return the first three elements in the first row with:

```
school_loc[1, c(1,2,3)]
```

```

##           X      Y OBJECTID
## 1 -86.5685 34.78337      1

```

R will return the values of `school_loc` that are in both the first row and the first, second, and third columns. Note that R won't actually remove these values from the data frame. R will give you a new set of values which are copies of the original values. You can then save this new set to an R object with R's assignment operator:

```

newdf <- school_loc[1, c(1,2,3)]
# newdf = school_loc[1, c(1,2,3)] # You can do this, but avoid.
head(newdf)

```

```

##           X      Y OBJECTID
## 1 -86.5685 34.78337      1

```

If you repeat a number in your index, R will return the corresponding value(s) more than once in your "subset." This code will return the first row of `school_loc` twice:

```
school_loc[c(1,1), c(1,2,2,3)]
```

```

##           X      Y      Y.1 OBJECTID
## 1 -86.5685 34.78337 34.78337      1
## 1.1 -86.5685 34.78337 34.78337      1

```

R's notation system is not limited to data frames. "You can use the same syntax to select values in any R object," as long as you supply one index for each dimension of the object. So, for example, you can subset a vector (which has one dimension) with a single index:

```

vec <- c(2,4,6,8,10)
print(vec[1:3])

```

```
## [1] 2 4 6
```

If you select two or more columns from a data frame, R will return a new data frame:

```
school_loc[1:2,1:2]
```

```

##           X      Y
## 1 -86.56850 34.78337

```

```
## 2 -86.79935 33.50570
```

However, if you select a single column, R will return a vector:

```
school_loc[1:2, 1]
```

```
## [1] -86.56850 -86.79935
```

If you would prefer a data frame instead, you can add the optional argument `drop = FALSE` between the brackets. This method also works for selecting a single column from a matrix or an array.

```
school_loc[1:2, 1, drop = FALSE]
```

```
##          X  
## 1 -86.56850  
## 2 -86.79935  
# school_loc[1:2, 1, drop = T]
```

3.1.2 Negative integers

Negative integers do the exact opposite of positive integers when indexing. R will return every element *except the elements in a negative index*. For example, `school_loc[-1, 1:3]` will return everything but the first row. `school_loc[-(2:52), 1:3]` will return the first row (and exclude everything else):

```
school_loc[-1, 1:3]  
school_loc[-(2:52), 1:3]  
school_loc[c(-1, -3), 1, drop=F]
```

Negative integers are a more efficient way to subset than positive integers if you want to include the majority of a data frame's rows or columns.

R will return an error if you try to pair a negative integer with a positive integer in the same index:

```
school_loc[c(-1, 1), 1]  
# Error in xj[i] : only 0's may be mixed with negative subscripts
```

However, you can use both negative and positive integers to subset an object if you use them in different indexes (e.g., if you use one in the rows index and one in the columns index, like `school_loc[-1, 1]`):

```
school_loc[-1, 1, drop=F]
```

3.1.3 Zero

What would happen if you used zero as an index? Zero is neither a positive integer nor a negative integer, but R will still use it to do a type of subsetting. R will return nothing from a dimension when you use zero as an index. This “creates an empty object”:

```
# data frame with 0 columns and 0 rows  
x <- school_loc[0, 0, drop = F]  
dim(x) # 0 0  
  
## [1] 0 0  
# creating same class, attributes, structure of objects  
# class(school_loc)  
# class(x)  
# attributes(school_loc)  
# attributes(x)
```

To be honest, indexing with zero is not very helpful and you will rarely need it.

3.1.4 Blank spaces

You can use a blank space to tell R to extract every value in a dimension. This lets you subset an object on one dimension but not the others, which is useful for extracting entire rows or columns from a data frame:

```
# This extracts the first row  
school_loc[1, ]  
# This extracts the first column  
school_loc[, 1, drop=F]
```

3.1.5 Logical values

If you supply a vector of TRUEs and FALSEs as your index, R will match each TRUE and FALSE to a row in your data frame (or a column depending on where you place the index). R will then return each row that corresponds to a TRUE.

It may help to imagine R reading through the data frame and asking, “Should I return the `_i_`th row of the data structure?” and then consulting the `_i_`th value of the index for its answer. For this system to work, your vector must be as long as the dimension you are trying to subset:

```
dim(school_loc) # [1] 7012 26  
  
## [1] 7012 26  
# You can use the multiplication symbol here as a shortcut to the FALSEs.  
  
num_cols <- ncol(school_loc) # 26  
head(school_loc)
```

```
##          X      Y OBJECTID UNITID             NAME  
## 1 -86.56850 34.78337    1 100654 Alabama A & M University  
## 2 -86.79935 33.50570    2 100663 University of Alabama at Birmingham  
## 3 -86.17401 32.36261    3 100690 Amridge University  
## 4 -86.64045 34.72456    4 100706 University of Alabama in Huntsville  
## 5 -86.29568 32.36432    5 100724 Alabama State University  
## 6 -87.52959 33.20702    6 100733 University of Alabama System Office  
##  
##            STREET      CITY STATE        ZIP STFIP CNTY  
## 1       4900 Meridian Street     Normal   AL 35762 01 01089  
## 2 Administration Bldg Suite 1070 Birmingham     AL 35294-0110 01 01073  
## 3           1200 Taylor Rd Montgomery     AL 36117-3553 01 01101  
## 4           301 Sparkman Dr Huntsville     AL 35899 01 01089  
## 5           915 S Jackson Street Montgomery     AL 36104-0271 01 01101  
## 6       500 University Blvd. East Tuscaloosa     AL 35401 01 01125  
##  
##          NM_CNTY LOCALE      LAT      LON CBSA             NM_CBSA  
## 1 Madison County     12 34.78337 -86.56850 26620 Huntsville, AL  
## 2 Jefferson County    12 33.50570 -86.79935 13820 Birmingham-Hoover, AL  
## 3 Montgomery County   12 32.36261 -86.17401 33860 Montgomery, AL  
## 4 Madison County     12 34.72456 -86.64045 26620 Huntsville, AL  
## 5 Montgomery County   12 32.36432 -86.29568 33860 Montgomery, AL  
## 6 Tuscaloosa County    12 33.20701 -87.52959 46220 Tuscaloosa, AL  
##  
##          CBSA_TYPE CSA             NMCSA NECTA NMNECTA CD SLDL  
## 1           1 290 Huntsville-Decatur, AL      N      N 0105 01019  
## 2           1 142 Birmingham-Hoover-Talladega, AL      N      N 0107 01055  
## 3           1 388 Montgomery-Selma-Alexander City, AL      N      N 0102 01074  
## 4           1 290 Huntsville-Decatur, AL      N      N 0105 01006  
## 5           1 388 Montgomery-Selma-Alexander City, AL      N      N 0107 01077  
## 6           1      N                           N      N 0107 01063
```

```

##      SLDU SCHOOLYEAR
## 1 01007 2020-2021
## 2 01018 2020-2021
## 3 01025 2020-2021
## 4 01002 2020-2021
## 5 01026 2020-2021
## 6 01021 2020-2021

school_loc[1, c(TRUE, FALSE, TRUE, rep(FALSE, 23))]

##          X OBJECTID
## 1 -86.5685      1
# View(school_loc)

```

This system may seem odd—who wants to type so many TRUEs and FALSEs? But it will become very powerful in the next section on modifying values.

3.1.6 Names

Finally, you can ask for the elements you want by name—if your object has names. This is a common way to extract the columns of a data frame, since columns almost always have names:

```

# If you do not remember the exact names of each column,
# you can first use the function colnames() to list all the column names.
colnames(school_loc)

```

```

## [1] "X"          "Y"          "OBJECTID"    "UNITID"     "NAME"
## [6] "STREET"     "CITY"        "STATE"       "ZIP"        "STFIP"
## [11] "CNTY"       "NMCNTY"     "LOCALE"      "LAT"        "LON"
## [16] "CBSA"       "NMCBSA"     "CBSATYPE"   "CSA"        "NMCSA"
## [21] "NECTA"      "NMNECTA"    "CD"         "SLDL"       "SLDU"
## [26] "SCHOOLYEAR"

school_loc[1:10, c("X", "Y", "NAME", "SCHOOLYEAR")]

```

	X	Y	NAME	SCHOOLYEAR
## 1	-86.56850	34.78337	Alabama A & M University	2020-2021
## 2	-86.79935	33.50570	University of Alabama at Birmingham	2020-2021
## 3	-86.17401	32.36261	Amridge University	2020-2021
## 4	-86.64045	34.72456	University of Alabama in Huntsville	2020-2021
## 5	-86.29568	32.36432	Alabama State University	2020-2021
## 6	-87.52959	33.20702	University of Alabama System Office	2020-2021
## 7	-87.54598	33.21188	The University of Alabama	2020-2021
## 8	-85.94527	32.92478	Central Alabama Community College	2020-2021
## 9	-86.96470	34.80679	Athens State University	2020-2021
## 10	-86.17754	32.36736	Auburn University at Montgomery	2020-2021

3.2 In-class exercises:

Now that you know the basics of R’s notation system, let’s put it to use.

1. Extract all values from the first ten rows from school_loc
2. Extract the 100th row from school_loc
3. Extract the 200th to 300th rows from school_loc, but only from the following columns: NAME, STREET, CITY, STATE, ZIP.

```
school_loc[1:10,]
```

	X	Y	OBJECTID	UNITID	NAME
## 1	-86.56850	34.78337	1	100654	Alabama A & M University
## 2	-86.79935	33.50570	2	100663	University of Alabama at Birmingham
## 3	-86.17401	32.36261	3	100690	Amridge University
## 4	-86.64045	34.72456	4	100706	University of Alabama in Huntsville
## 5	-86.29568	32.36432	5	100724	Alabama State University
## 6	-87.52959	33.20702	6	100733	University of Alabama System Office
## 7	-87.54598	33.21188	7	100751	The University of Alabama
## 8	-85.94527	32.92478	8	100760	Central Alabama Community College
## 9	-86.96470	34.80679	9	100812	Athens State University
## 10	-86.17754	32.36736	10	100830	Auburn University at Montgomery
	STREET		CITY	STATE	ZIP STFIP CNTY
## 1	4900 Meridian Street		Normal	AL	35762 01 01089
## 2	Administration Bldg Suite 1070		Birmingham	AL	35294-0110 01 01073
## 3	1200 Taylor Rd		Montgomery	AL	36117-3553 01 01101
## 4	301 Sparkman Dr		Huntsville	AL	35899 01 01089
## 5	915 S Jackson Street		Montgomery	AL	36104-0271 01 01101
## 6	500 University Blvd. East		Tuscaloosa	AL	35401 01 01125
## 7	739 University Blvd		Tuscaloosa	AL	35487-0100 01 01125
## 8	1675 Cherokee Rd	Alexander City		AL	35010 01 01123
## 9	300 N Beaty St		Athens	AL	35611 01 01083
## 10	7440 East Drive		Montgomery	AL	36117-3596 01 01101
	NMCNTY	LOCALE	LAT	LON	CBSA NMCSA NNECTA NMNNECTA CD SLDL
## 1	Madison County	12 34.78337	-86.56850	26620	Huntsville, AL
## 2	Jefferson County	12 33.50570	-86.79935	13820	Birmingham-Hoover, AL
## 3	Montgomery County	12 32.36261	-86.17401	33860	Montgomery, AL
## 4	Madison County	12 34.72456	-86.64045	26620	Huntsville, AL
## 5	Montgomery County	12 32.36432	-86.29568	33860	Montgomery, AL
## 6	Tuscaloosa County	12 33.20701	-87.52959	46220	Tuscaloosa, AL
## 7	Tuscaloosa County	12 33.21187	-87.54598	46220	Tuscaloosa, AL
## 8	Tallapoosa County	32 32.92478	-85.94527	10760	Alexander City, AL
## 9	Limestone County	31 34.80679	-86.96470	26620	Huntsville, AL
## 10	Montgomery County	12 32.36736	-86.17754	33860	Montgomery, AL
	CBSATYPE	CSA			
## 1	1	290	Huntsville-Decatur, AL	N	N 0105 01019
## 2	1	142	Birmingham-Hoover-Talladega, AL	N	N 0107 01055
## 3	1	388	Montgomery-Selma-Alexander City, AL	N	N 0102 01074
## 4	1	290	Huntsville-Decatur, AL	N	N 0105 01006
## 5	1	388	Montgomery-Selma-Alexander City, AL	N	N 0107 01077
## 6	1	N		N	N 0107 01063
## 7	1	N		N	N 0107 01063
## 8	2	388	Montgomery-Selma-Alexander City, AL	N	N 0103 01081
## 9	1	290	Huntsville-Decatur, AL	N	N 0105 01005
## 10	1	388	Montgomery-Selma-Alexander City, AL	N	N 0102 01074
	SLDU	SCHOOLYEAR			
## 1	01007	2020-2021			
## 2	01018	2020-2021			
## 3	01025	2020-2021			
## 4	01002	2020-2021			
## 5	01026	2020-2021			
## 6	01021	2020-2021			
## 7	01021	2020-2021			

```

## 8 01027 2020-2021
## 9 01001 2020-2021
## 10 01025 2020-2021

```

```
school_loc[c(1:10),]
```

	X	Y	OBJECTID	UNITID	NAME
## 1	-86.56850	34.78337	1	100654	Alabama A & M University
## 2	-86.79935	33.50570	2	100663	University of Alabama at Birmingham
## 3	-86.17401	32.36261	3	100690	Amridge University
## 4	-86.64045	34.72456	4	100706	University of Alabama in Huntsville
## 5	-86.29568	32.36432	5	100724	Alabama State University
## 6	-87.52959	33.20702	6	100733	University of Alabama System Office
## 7	-87.54598	33.21188	7	100751	The University of Alabama
## 8	-85.94527	32.92478	8	100760	Central Alabama Community College
## 9	-86.96470	34.80679	9	100812	Athens State University
## 10	-86.17754	32.36736	10	100830	Auburn University at Montgomery
			STREET	CITY STATE	ZIP STFIP CNTY
## 1	4900 Meridian Street		Normal	AL	35762 01 01089
## 2	Administration Bldg Suite 1070		Birmingham	AL	35294-0110 01 01073
## 3	1200 Taylor Rd		Montgomery	AL	36117-3553 01 01101
## 4	301 Sparkman Dr		Huntsville	AL	35899 01 01089
## 5	915 S Jackson Street		Montgomery	AL	36104-0271 01 01101
## 6	500 University Blvd. East		Tuscaloosa	AL	35401 01 01125
## 7	739 University Blvd		Tuscaloosa	AL	35487-0100 01 01125
## 8	1675 Cherokee Rd	Alexander City	AL	35010 01 01123	
## 9	300 N Beaty St	Athens	AL	35611 01 01083	
## 10	7440 East Drive	Montgomery	AL	36117-3596 01 01101	
	NMCNTY	LOCALE	LAT	LON	CBSA NMCBSA
## 1	Madison County	12 34.78337	-86.56850	26620	Huntsville, AL
## 2	Jefferson County	12 33.50570	-86.79935	13820	Birmingham-Hoover, AL
## 3	Montgomery County	12 32.36261	-86.17401	33860	Montgomery, AL
## 4	Madison County	12 34.72456	-86.64045	26620	Huntsville, AL
## 5	Montgomery County	12 32.36432	-86.29568	33860	Montgomery, AL
## 6	Tuscaloosa County	12 33.20701	-87.52959	46220	Tuscaloosa, AL
## 7	Tuscaloosa County	12 33.21187	-87.54598	46220	Tuscaloosa, AL
## 8	Tallapoosa County	32 32.92478	-85.94527	10760	Alexander City, AL
## 9	Limestone County	31 34.80679	-86.96470	26620	Huntsville, AL
## 10	Montgomery County	12 32.36736	-86.17754	33860	Montgomery, AL
	CBSATYPE	CSA		NMCSA	NECTA NMNECTA CD SLDL
## 1	1	290	Huntsville-Decatur, AL	N	N 0105 01019
## 2	1	142	Birmingham-Hoover-Talladega, AL	N	N 0107 01055
## 3	1	388	Montgomery-Selma-Alexander City, AL	N	N 0102 01074
## 4	1	290	Huntsville-Decatur, AL	N	N 0105 01006
## 5	1	388	Montgomery-Selma-Alexander City, AL	N	N 0107 01077
## 6	1	N		N	N 0107 01063
## 7	1	N		N	N 0107 01063
## 8	2	388	Montgomery-Selma-Alexander City, AL	N	N 0103 01081
## 9	1	290	Huntsville-Decatur, AL	N	N 0105 01005
## 10	1	388	Montgomery-Selma-Alexander City, AL	N	N 0102 01074
	SLDU	SCHOOLYEAR			
## 1	01007	2020-2021			
## 2	01018	2020-2021			
## 3	01025	2020-2021			
## 4	01002	2020-2021			

```

## 5 01026 2020-2021
## 6 01021 2020-2021
## 7 01021 2020-2021
## 8 01027 2020-2021
## 9 01001 2020-2021
## 10 01025 2020-2021

```

3.3 Dollar Signs and Double Brackets

Two types of object in R obey an optional second system of notation. You can extract values from data frames and lists with the \$ syntax. You will encounter the \$ syntax again and again as an R programmer, so let's examine how it works.

To select a column from a data frame, write the data frame's name and the column name separated by a \$. Notice that no quotes should go around the column name:

```

# For illustration purpose, I will just use a subset of the original school location dataset.
sub_school_loc <- school_loc[1:10,]
head(sub_school_loc)

```

```

##          X      Y OBJECTID UNITID           NAME
## 1 -86.56850 34.78337     1 100654 Alabama A & M University
## 2 -86.79935 33.50570     2 100663 University of Alabama at Birmingham
## 3 -86.17401 32.36261     3 100690 Amridge University
## 4 -86.64045 34.72456     4 100706 University of Alabama in Huntsville
## 5 -86.29568 32.36432     5 100724 Alabama State University
## 6 -87.52959 33.20702     6 100733 University of Alabama System Office
##             STREET        CITY STATE      ZIP STFIP CNTY
## 1       4900 Meridian Street    Normal   AL    35762 01 01089
## 2 Administration Bldg Suite 1070 Birmingham    AL 35294-0110 01 01073
## 3            1200 Taylor Rd Montgomery    AL 36117-3553 01 01101
## 4            301 Sparkman Dr Huntsville    AL    35899 01 01089
## 5            915 S Jackson Street Montgomery    AL 36104-0271 01 01101
## 6      500 University Blvd. East Tuscaloosa    AL    35401 01 01125
##          NMCTY LOCALE      LAT      LON CBSA           NMCSA
## 1    Madison County    12 34.78337 -86.56850 26620 Huntsville, AL
## 2 Jefferson County    12 33.50570 -86.79935 13820 Birmingham-Hoover, AL
## 3 Montgomery County    12 32.36261 -86.17401 33860 Montgomery, AL
## 4    Madison County    12 34.72456 -86.64045 26620 Huntsville, AL
## 5 Montgomery County    12 32.36432 -86.29568 33860 Montgomery, AL
## 6 Tuscaloosa County    12 33.20701 -87.52959 46220 Tuscaloosa, AL
##          CBSATYPE CSA           NMCSA NECTA NMNECTA CD SLDL
## 1          1 290 Huntsville-Decatur, AL      N      N 0105 01019
## 2          1 142 Birmingham-Hoover-Talladega, AL      N      N 0107 01055
## 3          1 388 Montgomery-Selma-Alexander City, AL      N      N 0102 01074
## 4          1 290 Huntsville-Decatur, AL      N      N 0105 01006
## 5          1 388 Montgomery-Selma-Alexander City, AL      N      N 0107 01077
## 6          1      N                           N      N 0107 01063
##          SLDU SCHOOLYEAR
## 1 01007 2020-2021
## 2 01018 2020-2021
## 3 01025 2020-2021
## 4 01002 2020-2021
## 5 01026 2020-2021
## 6 01021 2020-2021

```

```

sub_school_loc$NAME

## [1] "Alabama A & M University"           "University of Alabama at Birmingham"
## [3] "Amridge University"                  "University of Alabama in Huntsville"
## [5] "Alabama State University"            "University of Alabama System Office"
## [7] "The University of Alabama"            "Central Alabama Community College"
## [9] "Athens State University"             "Auburn University at Montgomery"

```

R will return all of the values in the column as a vector. This \$ notation is incredibly useful because you will often store the variables of your data sets as columns in a data frame. From time to time, you'll want to run a function like mean or median on the values in a variable. In R, these functions expect a vector of values as input:

```
mean(school_loc$X)
```

```
## [1] -90.35801
```

```
median(school_loc$Y)
```

```
## [1] 38.54617
```

```
max(school_loc$Y)
```

```
## [1] 71.3247
```

You can use the same \$ notation with the elements of a list, if they have names. This notation has an advantage with lists, too. If you subset a list in the usual way, R will return a new list that has the elements you requested. This is true even if you only request a single element. For example:

```
l <- list(numbers = c(1, 2), logical = TRUE, strings = c("a", "b", "c"))
```

```
1
```

```
## $numbers
```

```
## [1] 1 2
```

```
##
```

```
## $logical
```

```
## [1] TRUE
```

```
##
```

```
## $strings
```

```
## [1] "a" "b" "c"
```

And then subset it:

```
l[[1]] # access by index
```

```
## [1] 1 2
```

```
l$numbers
```

```
## [1] 1 2
```

```
l[1]
```

```
## $numbers
```

```
## [1] 1 2
```

typeof(l[1]) # list

typeof(l\$numbers) # double

The result is a smaller list with one element. That element is the vector `c(1, 2)`. This can be annoying because many R functions do not work with lists. For example, `sum(l[1])` will return an error. It would be horrible if once you stored a vector in a list, you could only ever get it back as a list:

```
sum(l[1])
# Error in sum(l[1]) : invalid 'type' (list) of argument
```

When you use the \$ notation, R will return the selected values as they are, with no list structure around them:

```
l$numbers

## [1] 1 2

## $numbers
## [1] 1 2

l$stringss

## [1] "a" "b" "c"
```

You can then immediately feed the results to a function:

```
sum(l$numbers)

## [1] 3
```

If the elements in your list do not have names (or you do not wish to use the names), you can use two brackets, instead of one, to subset the list. This notation will do the same thing as the \$ notation:

```
sum(l[[1]])

## [1] 3
```

In other words, if you subset a list with single-bracket notation, R will return a smaller list. If you subset a list with double-bracket notation, R will return just the values that were inside an element of the list. You can combine this feature with any of R's indexing methods:

```
l

## $numbers
## [1] 1 2
##
## $logical
## [1] TRUE
##
## $strings
## [1] "a" "b" "c"

l["numbers"]

## $numbers
## [1] 1 2
typeof(l["numbers"])

## [1] "list"
l[["numbers"]]

## [1] 1 2
typeof(l[["numbers"]])

## [1] "double"
```

```
sum(l[["numbers"]])
```

```
## [1] 3
```

This difference is subtle but important. In the R community, there is a popular, and helpful, way to think about it. Imagine that each list is a train and each element is a train car. When you use single brackets, R selects individual train cars and returns them as a new train. Each car keeps its contents, but those contents are still inside a train car (i.e., a list). When you use double brackets, R actually unloads the car and gives you back the contents.

4 Modifying Values

In this section, you will learn how to change the actual values that are stored inside your data frame. This is all adding up to something special: complete control of your data. You can now store your data in your computer, retrieve individual values at will, and use your computer to perform correct calculations with those values.

4.1 Changing Values in Place

You can use R's notation system to modify values within an R object. First, describe the value (or values) you wish to modify. Then use the assignment operator <- to overwrite those values. R will update the selected values in the original object. Let's put this into action with a real example:

```
vec <- c(1:10)
```

```
vec
```

```
## [1] 1 2 3 4 5 6 7 8 9 10
```

```
vec <- "text"
```

```
vec
```

```
## [1] "text"
```

Here's how you can select the first value of vec:

```
vec[1]
```

```
## [1] "text"
```

And here is how you can modify it:

```
vec
```

```
## [1] "text"
```

```
vec[1] <- 999
```

```
vec
```

```
## [1] "999"
```

```
vec[-1] <- 0
```

```
vec
```

```
## [1] "999"
```

You can replace multiple values at once as long as the number of new values equals the number of selected values:

```
vec <- 1:10
```

```
vec[c(1, 3, 5)] <- c(1000, 3000, 5000)
```

```
vec
```

```

## [1] 1000    2 3000    4 5000    6    7    8    9    10
vec[4:6] <- vec[4:6] + 1
vec

## [1] 1000    2 3000    5 5001    7    7    8    9    10

```

You can also create values that do not yet exist in your object. R will expand the object to accommodate the new values:

```

vec <- c(1:10)
vec

## [1] 1 2 3 4 5 6 7 8 9 10
length(vec)

## [1] 10
vec[12] <- 0
vec

## [1] 1 2 3 4 5 6 7 8 9 10 NA 0

```

Notice the NA in the vector. Our original vector has 10 elements, but here we added a new value at index 12 without specifying the value for the 11th element. R will automatically fill in the unspecified elements with NA. When writing your own code, try to avoid doing this. We usually try to minimize the amount of NAs in our dataframe.

Now you can add new variables to your data set:

```

colnames(school_loc)

## [1] "X"          "Y"          "OBJECTID"   "UNITID"     "NAME"
## [6] "STREET"     "CITY"        "STATE"       "ZIP"        "STFIP"
## [11] "CNTY"       "NMCNTY"     "LOCALE"      "LAT"        "LON"
## [16] "CBSA"       "NMCBSA"     "CBSATYPE"   "CSA"        "NMCSA"
## [21] "NECTA"      "NMNECTA"    "CD"         "SLDL"      "SLDU"
## [26] "SCHOOLYEAR"

school_loc$indices <- 1:7012
colnames(school_loc)

## [1] "X"          "Y"          "OBJECTID"   "UNITID"     "NAME"
## [6] "STREET"     "CITY"        "STATE"       "ZIP"        "STFIP"
## [11] "CNTY"       "NMCNTY"     "LOCALE"      "LAT"        "LON"
## [16] "CBSA"       "NMCBSA"     "CBSATYPE"   "CSA"        "NMCSA"
## [21] "NECTA"      "NMNECTA"    "CD"         "SLDL"      "SLDU"
## [26] "SCHOOLYEAR" "indices"

# View(school_loc)

```

You can also remove columns from a data frame (and elements from a list) by assigning them the symbol NULL:

```

school_loc$newcol <- 0
# View(school_loc)
school_loc$newcol <- NULL
colnames(school_loc)

## [1] "X"          "Y"          "OBJECTID"   "UNITID"     "NAME"
## [6] "STREET"     "CITY"        "STATE"       "ZIP"        "STFIP"

```

```

## [11] "CNTY"          "NMCNTY"        "LOCALE"        "LAT"           "LON"
## [16] "CBSA"          "NMCBSA"        "CBSATYPE"      "CSA"           "NMCSA"
## [21] "NECTA"         "NMNECTA"       "CD"            "SLDL"          "SLDU"
## [26] "SCHOOLYEAR"    "indices"

```

You can single out just the values by subsetting the columns dimension. You can also assign a new set of values to these old values. The set of new values will have to be the same size as the set of values that you are replacing.

```

# To illustrate subsetting values, we add the new variable back.
school_loc$newcol <- 1:7012
head(school_loc$newcol, 10)

```

```

## [1] 1 2 3 4 5 6 7 8 9 10
school_loc$newcol[c(1,3,5,7,9)] <- c(111, 333, 555, 777, 999)
head(school_loc$newcol, 10)

```

```

## [1] 111 2 333 4 555 6 777 8 999 10
# If you would like to set them to the same value, simply do:
school_loc$newcol[c(2,4,6,8,10)] <- 1000
head(school_loc$newcol, 10)

```

```

## [1] 111 1000 333 1000 555 1000 777 1000 999 1000

```

Notice that the values change in place. You do not end up with a new copy of the school_loc dataframe. The new values will appear inside the school_loc dataframe.

The same technique will work whether you store your data in a vector, matrix, list, or data frame. Just describe the values that you want to change with R's notation system, then assign over those values with R's assignment operator.

4.2 Logical Subsetting

Do you remember R's logical index system, *logicals*? To recap, you can select values with a vector of `TRUEs` and `FALSEs`. The vector must be the same length as the dimension that you wish to subset. R will return every element that matches a `TRUE`:

```

vec2 <- c(1,0,2,0,3,4,5,0,9)
vec2

## [1] 1 0 2 0 3 4 5 0 9
vec2[c(TRUE, FALSE, FALSE, FALSE, FALSE, FALSE, TRUE, FALSE, TRUE)] # filter

## [1] 1 5 9

```

At first glance, this system might seem impractical. Who wants to type out long vectors of `TRUEs` and `FALSEs`? No one. But you don't have to. You can let a logical test create a vector of `TRUEs` and `FALSEs` for you.

4.2.1 Logical Tests

A logical test is a comparison like “is one less than two?”, $1 < 2$, or “is three greater than four?”, $3 > 4$. R provides seven logical operators that you can use to make comparisons.

Each operator returns a `TRUE` or a `FALSE`. If you use an operator to compare vectors, R will do ‘element-wise comparisons’ — just like it does with the arithmetic operators:

```
1 < 2
```

```
## [1] TRUE
```

Operator	Syntax	Tests
>	a > b	Is a greater than b?
>=	a >= b	Is a greater than or equal to b?
<	a < b	Is a less than b?
<=	a <= b	Is a less than or equal to b?
==	a == b	Is a equal to b?
!=	a != b	Is a not equal to b?
%in%	a %in% c(a, b, c)	Is a in the group c(a, b, c)?

Figure 1: R's Logical Operators

```
1 > c(0, 1, 2)
## [1] TRUE FALSE FALSE
c(1, 2, 3) == c(3, 2, 1)
## [1] FALSE TRUE FALSE
```

%in% is the only operator that does ‘not’ do normal element-wise execution. %in% tests whether the value(s) on the left side are in the vector on the right side. If you provide a vector on the left side, %in% will not pair up the values on the left with the values on the right and then do element-wise tests. Instead, %in% will independently test whether each value on the left is somewhere in the vector on the right:

```
1 %in% c(3, 4, 5)
## [1] FALSE
c(1, 2) %in% c(3, 4, 5)
## [1] FALSE FALSE
c(1, 2, 3) %in% c(3, 4, 5)
## [1] FALSE FALSE TRUE
c(1, 2, 3, 4) %in% c(3, 4, 5)
## [1] FALSE FALSE TRUE TRUE
# works for strings
"a" %in% c("a", "b", "c")
## [1] TRUE
"A" %in% c("a") # case sensitive
## [1] FALSE
```

```
"a" %in% c("abc") # %in% checks exact matches
```

```
## [1] FALSE
```

Notice that you test for equality with a double equals sign, `==`, and not a single equals sign, `=`, which is another way to write the assignment operator `<-`. It is easy to forget and use `a = b` to test if `a` equals `b`. Unfortunately, you'll be in for a nasty surprise. R won't return a `TRUE` or `FALSE`, because it won't have to: `a` will equal `b`, because you just ran the equivalent of `a <- b` where you manually set the value of `a` to be `b`.

You can compare any two R objects with a logical operator; however, logical operators make the most sense if you compare two objects of the same data type. If you compare objects of different data types, R will use its coercion rules to coerce the objects to the same type before it makes the comparison and the result might be hard to interpret.

We again use the school location dataset as our example.

```
colnames(school_loc)
# First extract the column "X"
school_loc$X
# Use the < operator to check whether each value is less than -100
school_loc$X < -100
# Use the sum() function to check how many TRUEs there are
# TRUE = 1 / FALSE = 0
sum(school_loc$X < -100)

# What does this do?
school_loc$Y[school_loc$X < -100]
# uses the logical vector to index or subset the Y column.
# extract all values from the Y column where the corresponding values in the X column are TRUE(less than -100)

# What does this do?
school_loc$Y[school_loc$X < -100] <- 0
# View(school_loc)
```

To summarize, you can use a logical test to select values within an object.

Logical subsetting is a powerful technique because it lets you quickly identify, extract, and modify individual values in your data set. When you work with logical subsetting, you do not need to know where in your data set a value exists. You only need to know how to describe the value with a logical test.

Logical subsetting is one of the things R does best. In fact, logical subsetting is a key component of vectorized programming, a coding style that lets you write fast and efficient R code.

4.3 In-class exercises:

Let's practice with the dataset TFR.

1. Find the dimension of TFR. How many observations and variables?
2. Add a new column called index to TFR, where the values are the indices for the rows.
3. Select the first row of TFR.
4. Select the first column of TFR.
5. Get the last row of TFR.
6. How many values in the column TFR are greater than 8?
7. How many values in the column Childbearing are over 25?
8. What does this code do?

```
TFR$Childbearing[TFR$Childbearing > 25]
```

9. Set all values in the column ChildBearing that are less than 25 to 0.

```
TFR <- read.csv("dataset.csv")
sum(TFR$TFR > 8)

## [1] 67
nrow(TFR[TFR$TFR > 8,])

## [1] 67
typeof(TFR[TFR$TFR > 8,])

## [1] "list"
head(TFR[TFR$TFR > 8, ])

##          Country Uncode Year      TFR InfMRateCME InfMRateUN U5MRateCME U5MRateUN
## 5215      Iraq    368 1950 8.3980      246.2    280.6580     364.3   407.6411
## 5216      Iraq    368 1951 8.1784      231.0    269.5758     343.1   391.3986
## 5625      Jordan   400 1964 8.0340      87.1    100.3606     123.8   149.2374
## 5626      Jordan   400 1965 8.0570      82.8    95.5380     117.1   140.3478
## 5627      Jordan   400 1966 8.0530      78.9    91.9096     110.9   133.8606
## 5628      Jordan   400 1967 8.0330      75.2    88.2812     105.1   127.3735
##          LifeExpB MtoFBirth  MtoF04 Pop1564 Pop1564Female GDPpc GDPpcGrowth
## 5215 30.9550        1.065 1.078620 59.78318    60.41158      NA       NA
## 5216 32.3510        1.065 1.067612 59.09168    59.80392      NA       NA
## 5625 55.7927        1.054 1.079269 51.52381    51.79283 3744.551 0.1159344
## 5626 56.5522        1.054 1.073565 51.33929    51.86567 3991.914 0.0639687
## 5627 57.3033        1.054 1.069025 50.94963    51.03448 3890.935 -0.0256205
## 5628 58.0440        1.054 1.066867 50.90909    51.18483 4012.081 0.0306606
##          Yschooling YschoolF1549 GenrollPrim ChildBearing CountryCode
## 5215      0.240     0.1059529      NA     28.6970      IRQ
## 5216      0.258     0.1174685      NA     28.6970      IRQ
## 5625      2.882     1.8347658      NA     30.0230      JOR
## 5626      3.000     1.9325727      NA     30.0870      JOR
## 5627      3.094     2.0368810      NA     30.1594      JOR
## 5628      3.188     2.1411893      NA     30.2318      JOR
typeof(TFR$TFR > 8)

## [1] "logical"
nrow(TFR[TFR$TFR > 8,])

## [1] 67
```

4.3.1 Boolean Operators

Boolean operators are things like and (`&`) and or (`|`). They “collapse the results of multiple logical tests into a single TRUE or FALSE”. R has six boolean operators.

To use a Boolean operator, place it between two complete logical tests. R will execute each logical test and then use the Boolean operator to combine the results into a single TRUE or FALSE. If you do not supply a complete test to each side of the operator, R will return an error.

When used with vectors, Boolean operators will follow the same element-wise execution as arithmetic and logical operators:

Operator	Syntax	Tests
&	cond1 & cond2	Are both cond1 and cond2 true?
	cond1 cond2	Is one or more of cond1 and cond2 true?
xor	xor(cond1, cond2)	Is exactly one of cond1 and cond2 true?
!	!cond1	Is cond1 false? (e.g., ! flips the results of a logical test)
any	any(cond1, cond2, cond3, ...)	Are any of the conditions true?
all	all(cond1, cond2, cond3, ...)	Are all of the conditions true?

Figure 2: Boolean Operators

```

a <- c(1, 2, 3)
b <- c(1, 2, 3)
c <- c(1, 2, 4)
a == b

## [1] TRUE TRUE TRUE
b == c

## [1] TRUE TRUE FALSE
a == b & b == c

## [1] TRUE TRUE FALSE
all(a == b & b == c)

## [1] FALSE
any(a == b & b == c)

## [1] TRUE
# (TRUE TRUE TRUE) & (TRUE TRUE FALSE)

```

4.4 In-class exercises:

If you think you have the hang of logical tests, try converting these sentences into tests written with R code. To help you out, I've defined some R objects after the sentences that you can use to test your answers:

1. Is w positive?
2. Is x greater than 10 and less than 20?
3. Is object y the word February?
4. Is every value in z a day of the week?
5. What does this expression evaluate to and why?

```
w <- c(-1, 0, 1)
x <- c(5, 15)
y <- "February"
z <- c("Monday", "Tuesday", "Friday")
(TRUE + TRUE) * FALSE
```

```
## [1] 0
```

6. Use logical operators to output only those rows of data in TFR where column Year is between 1950 and 1955 inclusively.
7. Use logical operators to output only those rows of data where column TFR is equal to 7.45 and column Year is greater than 1960.
8. Use logical operators to output only the even rows of the dataframe.
9. Use logical operators and change every 4th element in column LifeExpB to 0.

```
data <- TFR
# 6
head(data[data$Year >= 1950 & data$Year <= 1955,])
```

```
##          Country Uncode Year    TFR InfMRateCME InfMRateUN U5MRateCME U5MRateUN
## 1 Afghanistan      4 1950 7.45        NA 304.5940        NA 438.0103
## 2 Afghanistan      4 1951 7.45        NA 299.6836        NA 431.6065
## 3 Afghanistan      4 1952 7.45        NA 294.7732        NA 425.2027
## 4 Afghanistan      4 1953 7.45        NA 289.8628        NA 418.7989
## 5 Afghanistan      4 1954 7.45        NA 284.9524        NA 412.3951
## 6 Afghanistan      4 1955 7.45        NA 280.0420        NA 405.9913
##   LifeExpB MtoFbirth    MtoF04 Pop1564 Pop1564Female GDPpc GDPpcGrowth
## 1 26.0690         1.06 0.9523352 56.08875     54.25678    NA      NA
## 2 26.5736         1.06 0.9524428 55.82908     54.06208    NA      NA
## 3 27.0782         1.06 0.9728808 55.74039     54.08136    NA      NA
## 4 27.5828         1.06 0.9952082 55.71038     54.13613    NA      NA
## 5 28.0874         1.06 1.0122050 55.71779     54.22245    NA      NA
## 6 28.5920         1.06 1.0222970 55.68319     54.21412    NA      NA
##   Yschooling YschoolF1549 GenrollPrim ChildBearing CountryCode
## 1      0.270  0.08679564        NA     29.835      AFG
## 2      0.278  0.08758149        NA     29.835      AFG
## 3      0.286  0.08836733        NA     29.835      AFG
## 4      0.294  0.08915317        NA     29.835      AFG
## 5      0.302  0.08993901        NA     29.835      AFG
## 6      0.310  0.09072485        NA     29.835      AFG
```

```
# 7
head(data[data$TFR == 7.45 & data$Year > 1960,])
```

```
##          Country Uncode Year    TFR InfMRateCME InfMRateUN U5MRateCME U5MRateUN
## 12 Afghanistan      4 1961 7.45        240.5 251.1132      356.5 368.4389
## 13 Afghanistan      4 1962 7.45        236.3 246.7364      350.6 362.9053
## 14 Afghanistan      4 1963 7.45        232.3 242.3596      345.0 357.3718
## 15 Afghanistan      4 1964 7.45        228.5 237.9828      339.7 351.8383
## 16 Afghanistan      4 1965 7.45        224.6 233.6060      334.1 346.3048
## 17 Afghanistan      4 1966 7.45        220.7 229.9888      328.7 341.2930
##   LifeExpB MtoFbirth    MtoF04 Pop1564 Pop1564Female GDPpc GDPpcGrowth
## 12 32.7774         1.06 1.017161 54.74086     53.69702    NA      NA
## 13 33.2199         1.06 1.016234 54.53767     53.62031    NA      NA
```

```

## 14 33.6579    1.06 1.021365 54.46916    53.74811    NA    NA
## 15 34.0929    1.06 1.027291 54.43051    53.94208    NA    NA
## 16 34.5254    1.06 1.030376 54.34323    53.93676    NA    NA
## 17 34.9574    1.06 1.033343 54.02503    53.72803    NA    NA
##   Yschooling YschoolF1549 GenrollPrim ChildBearing CountryCode
## 12      0.380    0.1088078      NA    29.835    AFG
## 13      0.390    0.1101457      NA    29.835    AFG
## 14      0.400    0.1114836      NA    29.835    AFG
## 15      0.410    0.1128215      NA    29.835    AFG
## 16      0.420    0.1141593      NA    29.835    AFG
## 17      0.474    0.1291685      NA    29.835    AFG
# 8
head(data[c(F,T),])

##          Country Uncode Year  TFR InfMRateCME InfMRateUN U5MRateCME U5MRateUN
## 2 Afghanistan     4 1951 7.45        NA 299.6836        NA 431.6065
## 4 Afghanistan     4 1953 7.45        NA 289.8628        NA 418.7989
## 6 Afghanistan     4 1955 7.45        NA 280.0420        NA 405.9913
## 8 Afghanistan     4 1957 7.45        NA 270.2212        NA 393.1837
## 10 Afghanistan    4 1959 7.45        NA 260.4004        NA 380.3761
## 12 Afghanistan    4 1961 7.45    240.5 251.1132    356.5 368.4389
## LifeExpB MtoFbirth MtoF04 Pop1564 Pop1564Female GDPpc GDPpcGrowth
## 2 26.5736       1.06 0.9524428 55.82908 54.06208    NA    NA
## 4 27.5828       1.06 0.9952082 55.71038 54.13613    NA    NA
## 6 28.5920       1.06 1.0222970 55.68319 54.21412    NA    NA
## 8 29.6012       1.06 1.0493490 55.25483 53.91750    NA    NA
## 10 30.6104      1.06 1.0313490 55.02604 53.76849    NA    NA
## 12 32.7774      1.06 1.0171610 54.74086 53.69702    NA    NA
##   Yschooling YschoolF1549 GenrollPrim ChildBearing CountryCode
## 2      0.278    0.08758149      NA    29.835    AFG
## 4      0.294    0.08915317      NA    29.835    AFG
## 6      0.310    0.09072485      NA    29.835    AFG
## 8      0.334    0.09742289      NA    29.835    AFG
## 10     0.358    0.10412092      NA    29.835    AFG
## 12     0.380    0.10880781      NA    29.835    AFG
# 9
data$LifeExpB[c(F,F,F,T)] <- 0

# %%

```

4.5 Missing Information (NA)

Missing information problems happen frequently in data science. Usually, they are more straightforward: you don't know a value because the measurement was lost, corrupted, or never taken to begin with. R has a way to help you manage these missing values.

The NA character is a special symbol in R. It stands for "not available" and can be used as a placeholder for missing information. R will treat NA exactly as you should want missing information treated. For example, what result would you expect if you add 1 to a piece of missing information?

```
1 + NA
```

```
## [1] NA
```

R will return a second piece of missing information. It would not be correct to say that $1 + \text{NA} = 1$ because

there is a good chance that the missing quantity is not zero. You do not have enough information to determine the result.

What if you tested whether a piece of missing information is equal to 1?

```
NA == 3.14
```

```
## [1] NA
```

Again, your answer would be something like “I do not know if this is equal to one,” that is, NA. Generally, NAs will propagate whenever you use them in an R operation or function. This can save you from making errors based on missing data.

4.5.1 Inside functions: `na.rm` argument

Missing values can help you work around holes in your data sets, but they can also create some frustrating problems. Suppose, for example, that you’ve collected 1,000 observations and wish to take their average with R’s mean function. If even one of the values is NA, your result will be NA:

```
c(NA, 1:50)
```

```
## [1] NA 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24
## [26] 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46 47 48 49
## [51] 50
```

```
mean(c(NA, 1:50))
```

```
## [1] NA
```

Understandably, you may prefer a different behavior. Most R functions come with the optional argument, `na.rm`, which stands for NA remove. R will ignore NAs when it evaluates a function if you add the argument `na.rm = TRUE`:

```
mean(c(NA, 1:50), na.rm = TRUE)
```

```
## [1] 25.5
```

```
mean(c(1:50)) # which returns the same value as above
```

```
## [1] 25.5
```

4.5.2 Test for NA: `is.na()`

On occasion, you may want to identify the NAs in your data set with a logical test, but that too creates a problem. How would you go about it? If something is a missing value, any logical test that uses it will return a missing value, even this test:

```
NA == NA
```

```
## [1] NA
```

Which means that tests like this won’t help you find missing values:

```
c(1, 2, 3, NA) == NA
```

```
## [1] NA NA NA NA
```

But don’t worry too hard; R supplies a special function that can test whether a value is an NA. The function is sensibly named `is.na()`:

```
is.na(NA)
```

```
## [1] TRUE
```

```

vec3 <- c(1, 2, 3, NA)
is.na(vec3)

## [1] FALSE FALSE FALSE TRUE
# What does this do?
school_loc$Y[school_loc$X < -100] <- NA
# View(school_loc)

```

4.5.3 Omit the Entire Rows: na.omit()

You can also remove all rows containing NA's in any column by using the `na.omit()` function. For example,

```

dim(school_loc)
summary(school_loc) # 7012 obs.
new_df <- na.omit(school_loc)
dim(new_df)
summary(new_df) # 5470 obs

```

4.6 In-class exercises:

1. What is the length of X?

```

X <- c (123,0,NA,8,NA,200)
na.omit(X)

```

```

## [1] 123   0    8  200
## attr(,"na.action")
## [1] 3 5
## attr(,"class")
## [1] "omit"
X

```

```

## [1] 123   0   NA   8   NA  200

```

2. Can you find all occurrences of NA in X? How can you find the total number of NAs in X?
3. Can you remove all occurrences of NA in X?
4. Can you replace all occurrences of NA with 88?
5. We will use the TFR dataset again. Create a new dataframe TFR2 where you drop all NA's.
6. Find the dimension of TFR2. How many observations are left?
7. Create a new dataframe TFR3 where you remove all rows with NA values in the GDPpc column.

```

TFR3 <- TFR[!is.na(TFR$GDPpc), ]
# View(TFR3)

```